

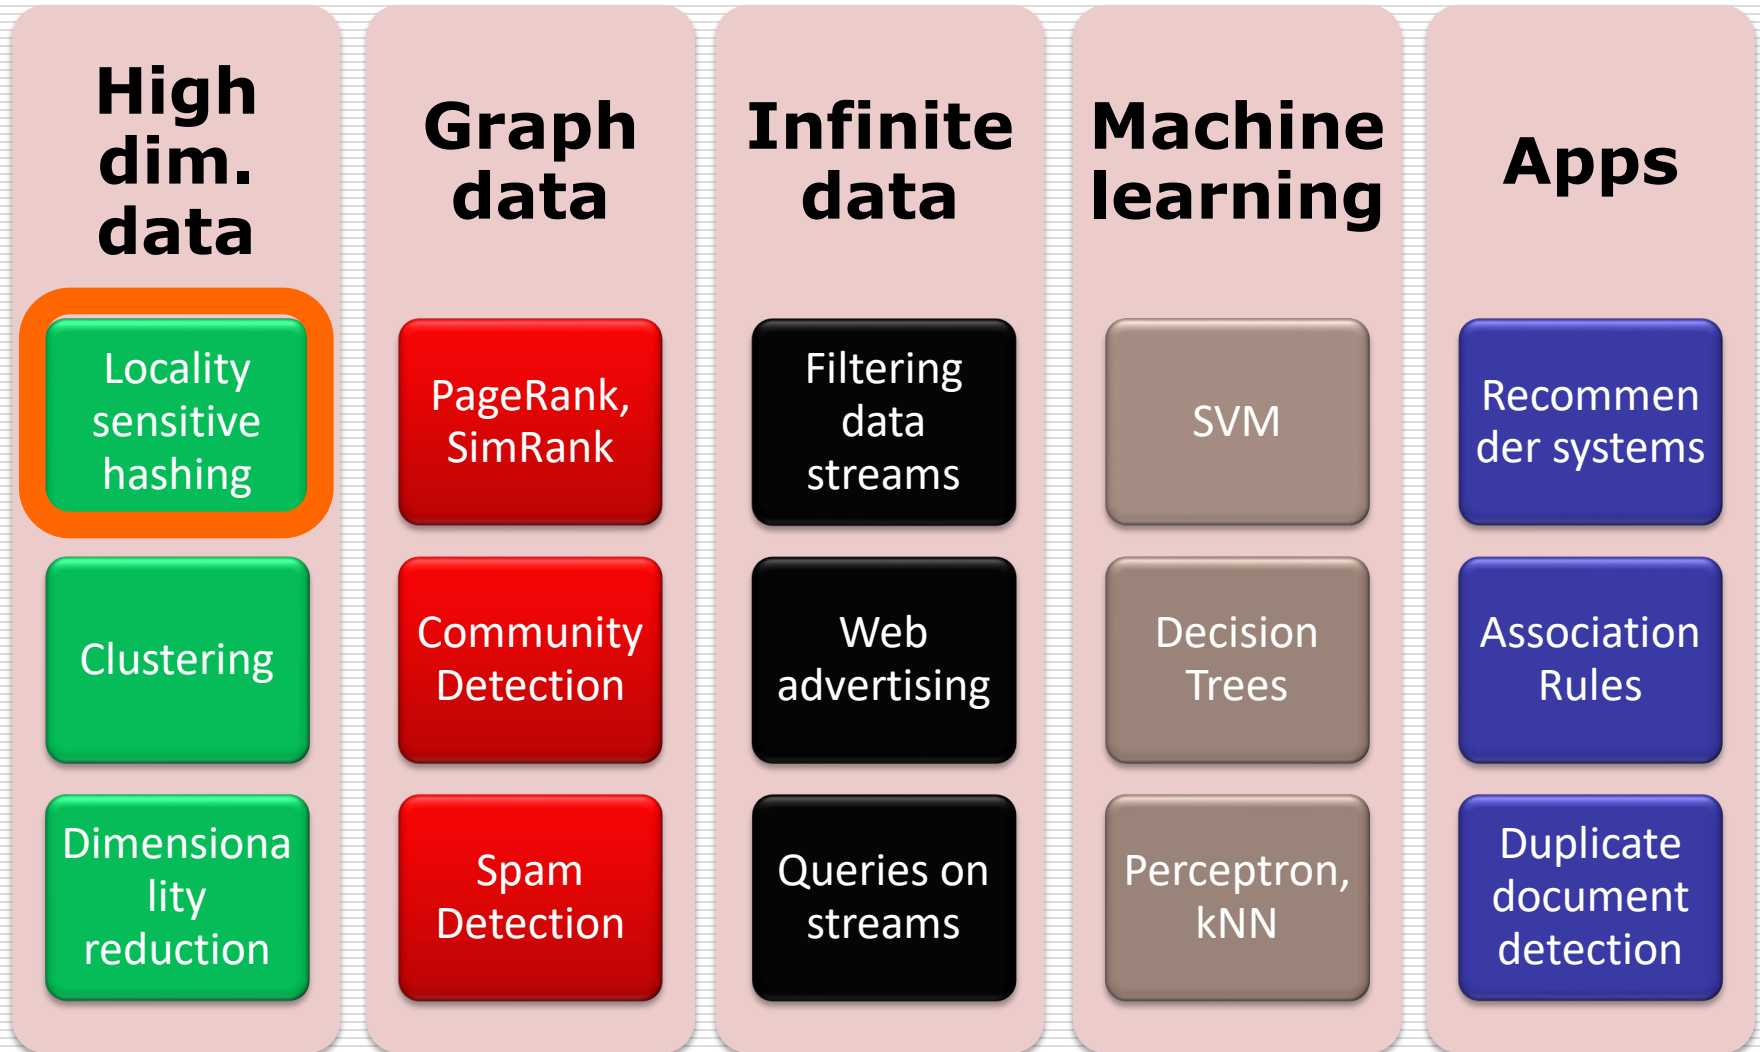
# 大数据计算及应用(三)

---

## Finding Similar Items: Locality Sensitive Hashing


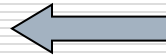
# Agenda

---



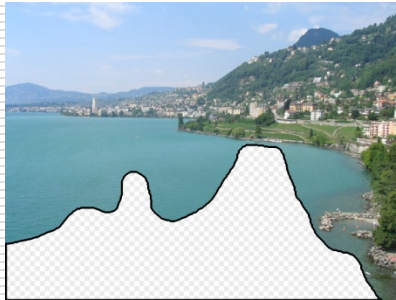
# Scene Completion Problem

---



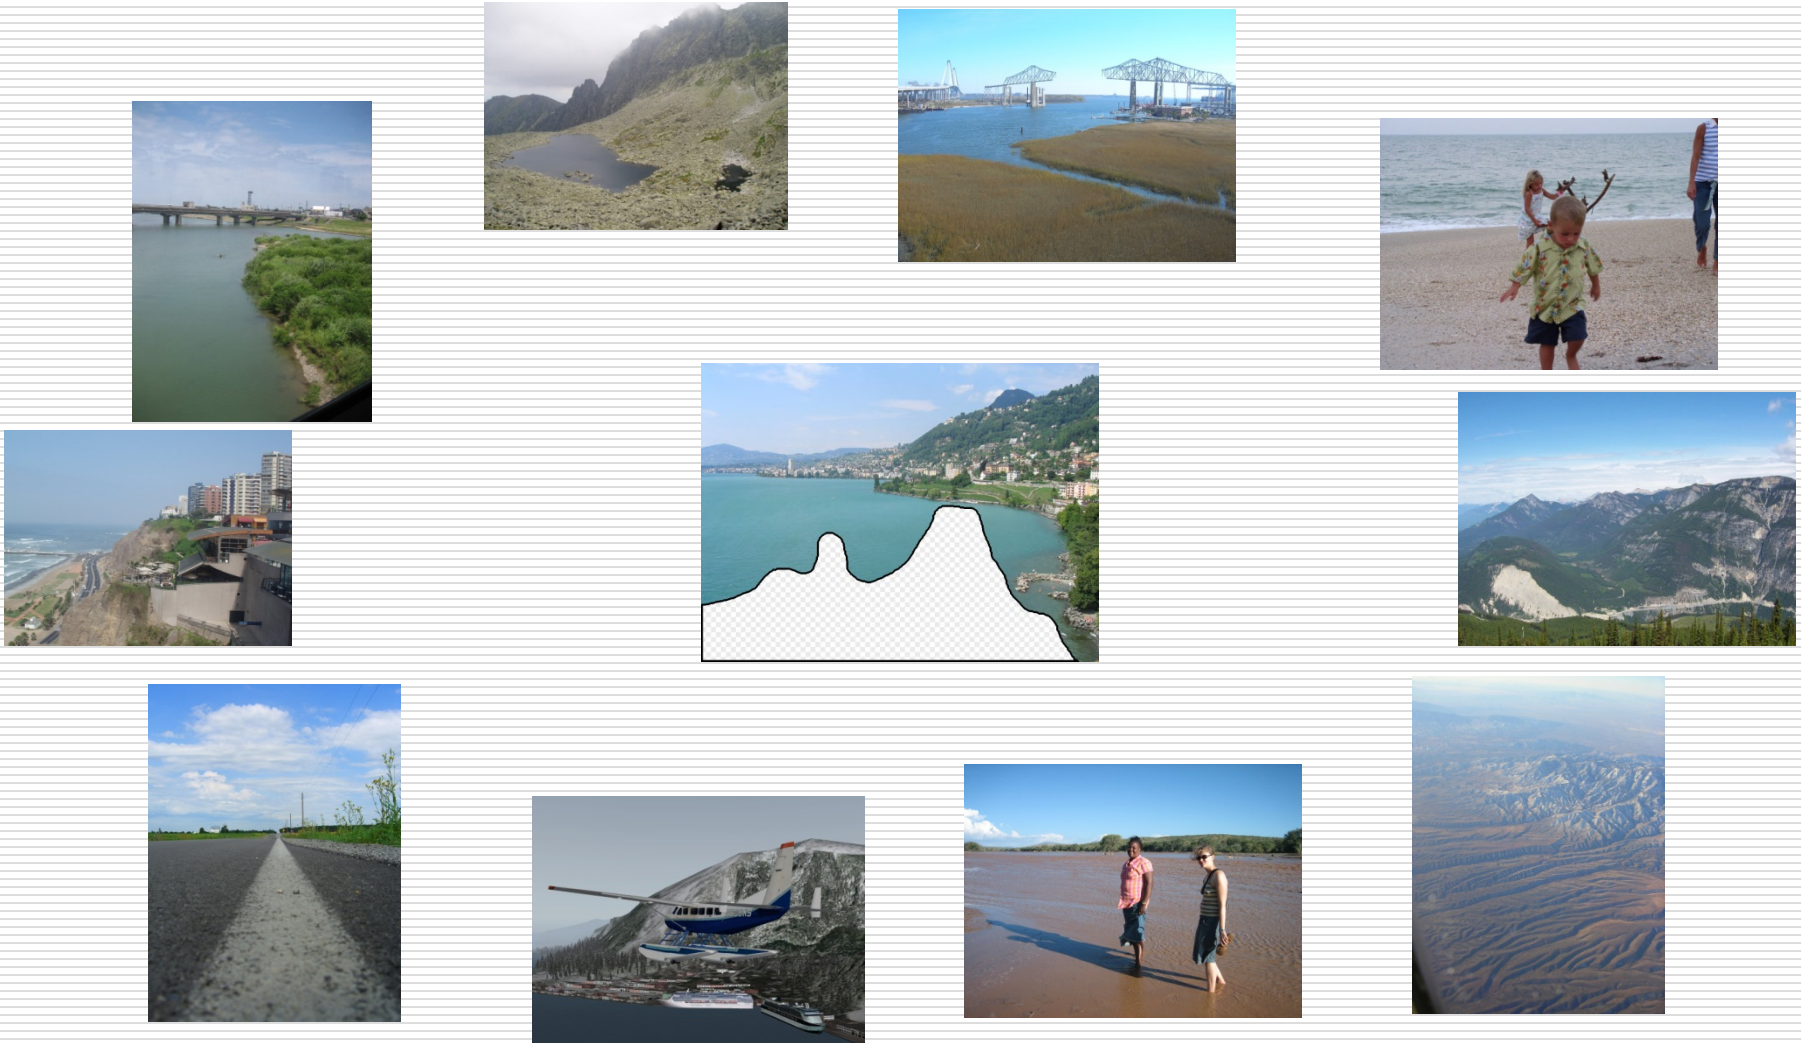
# Scene Completion Problem

---



# Scene Completion Problem

---



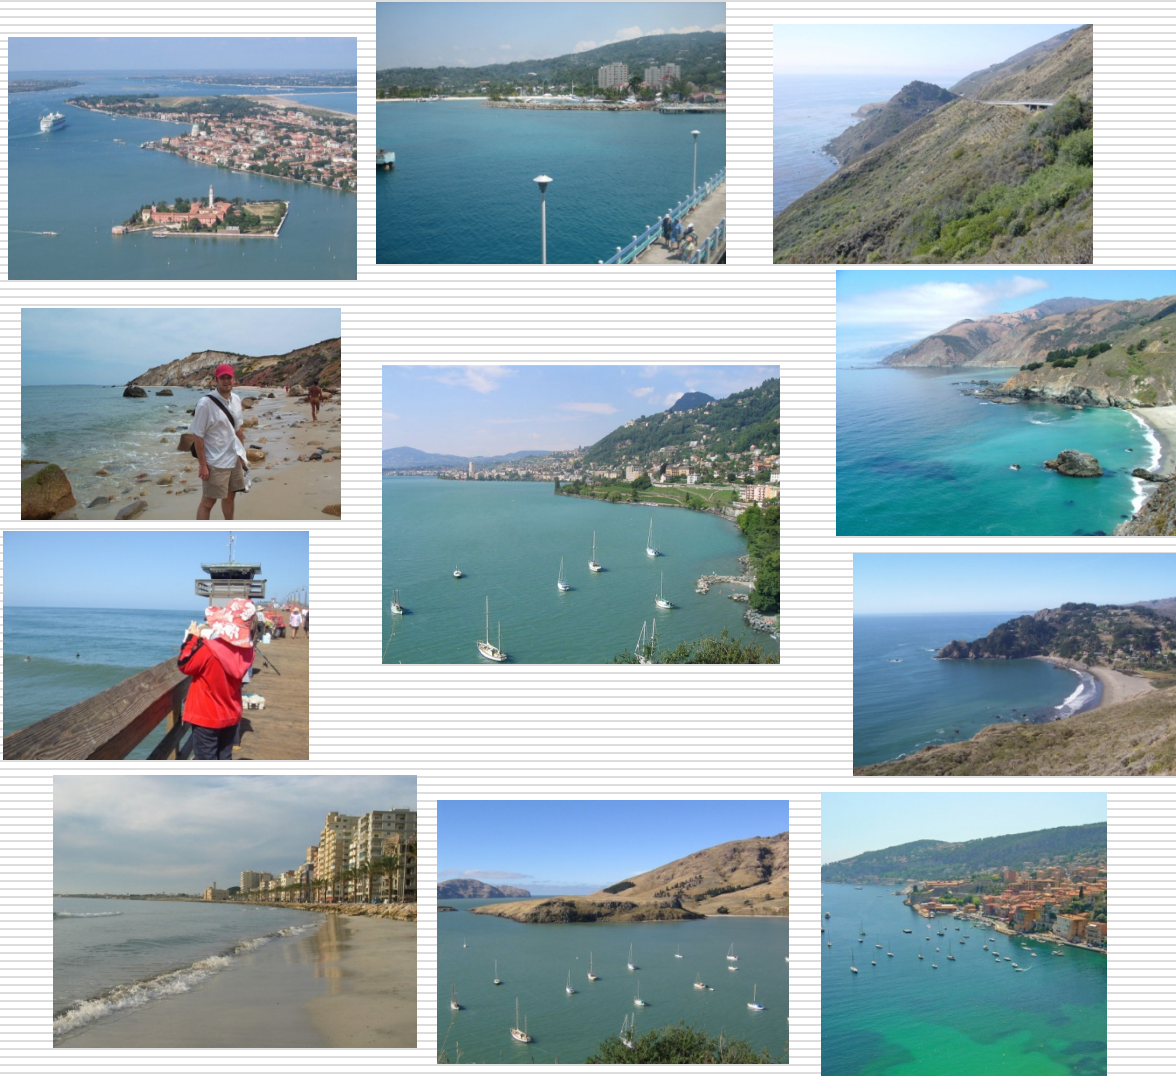
---

10 nearest neighbors from a collection of 2.3M images



# Scene Completion Problem

---



---

**10 nearest neighbors from a collection of 2.3M images**

# A Common Metaphor

---

- Many problems can be expressed as finding “similar” sets:
  - Find near-neighbors in high-dimensional space
- Examples:
  - Pages with similar words
    - For duplicate detection, classification by topic
  - Customers who purchased similar products
    - Products with similar customer sets
  - Images with similar features
    - Users who visited similar websites



# Problem for Today's Lecture

---

□ Given: High dimensional data points  $x_1, x_2, \dots$

■ For example: Image is a long vector of pixel colors

$$\begin{bmatrix} 1 & 0 & 1 \\ 0 & 0 & 1 \\ 0 & 1 & 0 \end{bmatrix} \rightarrow [1 \ 0 \ 1 \ 0 \ 0 \ 1 \ 0 \ 1 \ 0]$$

<i>img1</i>	<i>img2</i>
1	1
0	1
1	1
0	1
0	0
1	1
0	1
1	1
0	0

Is *img1* similar to *img2*?



# Problem for Today's Lecture

---

□ Given: High dimensional data points  $x_1, x_2, \dots$

■ Another example:


Document is a long vector of words existence

	<i>doc1</i>	<i>doc2</i>
angle	1	1
before	0	1
country	1	1
end	0	1
food	0	0
good	1	1
house	0	1
live	1	1
use	0	0

Is *doc1* similar to *doc2*?

# Problem for Today's Lecture

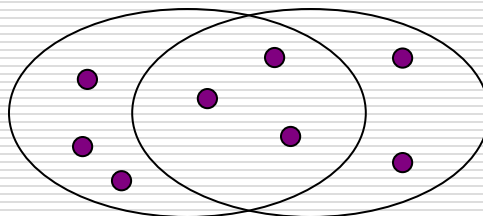
- Goal: Find **all pairs of documents**  $(x_i, x_j)$  that are within some distance threshold  $d(x_i, x_j) \leq s$



# Distance Measures

---

- **Goal: Find near-neighbors in high-dim. space**
  - We formally define “near neighbors” as points that are a “small distance” apart
- For each application, we first need to define what “distance” means
- **Today: Jaccard distance/similarity**
  - The **Jaccard similarity** of two **sets** is the size of their intersection divided by the size of their union:  
$$\text{sim}(C_1, C_2) = |C_1 \cap C_2| / |C_1 \cup C_2|$$
  - **Jaccard distance:**  $d(C_1, C_2) = 1 - |C_1 \cap C_2| / |C_1 \cup C_2|$



3 in intersection

8 in union

Jaccard similarity = 3/8

Jaccard distance = 5/8

# Encoding Sets as Bit Vectors

---

- Encode sets using 0/1 (bit, boolean) vectors
  - One dimension per element in the universal set
- Interpret set intersection as bitwise AND, and set union as bitwise OR
- Example:  $C_1 = 10111$ ;  $C_2 = 10011$ 
  - Size of intersection = 3; size of union = 4,
  - Jaccard similarity (not distance) =  $3/4$
  - Distance:  $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 1/4$

# Task: Finding Similar Documents

---

- **Goal:** Given a large number ( $N$  in the millions or billions) of documents, find “near duplicate” pairs
- **Applications:**
  - Mirror websites, or approximate mirrors
    - Don’t want to show both in search results
  - Similar news articles at many news sites
    - Cluster articles by “same story”
- **Problems:**
  - Many small pieces of one document can appear out of order in another
  - Documents are so large or so many that they cannot fit in main memory
  - Too many documents to compare all pairs



# 3 Essential Steps for Similar Docs

---

1. *Shingling*: Convert documents to sets
2. *Min-Hashing*: Convert large sets to short signatures, while preserving similarity
3. *Locality-Sensitive Hashing*: Focus on pairs of signatures likely to be from similar documents
  - Candidate pairs!

# Problem for Today's Lecture

---

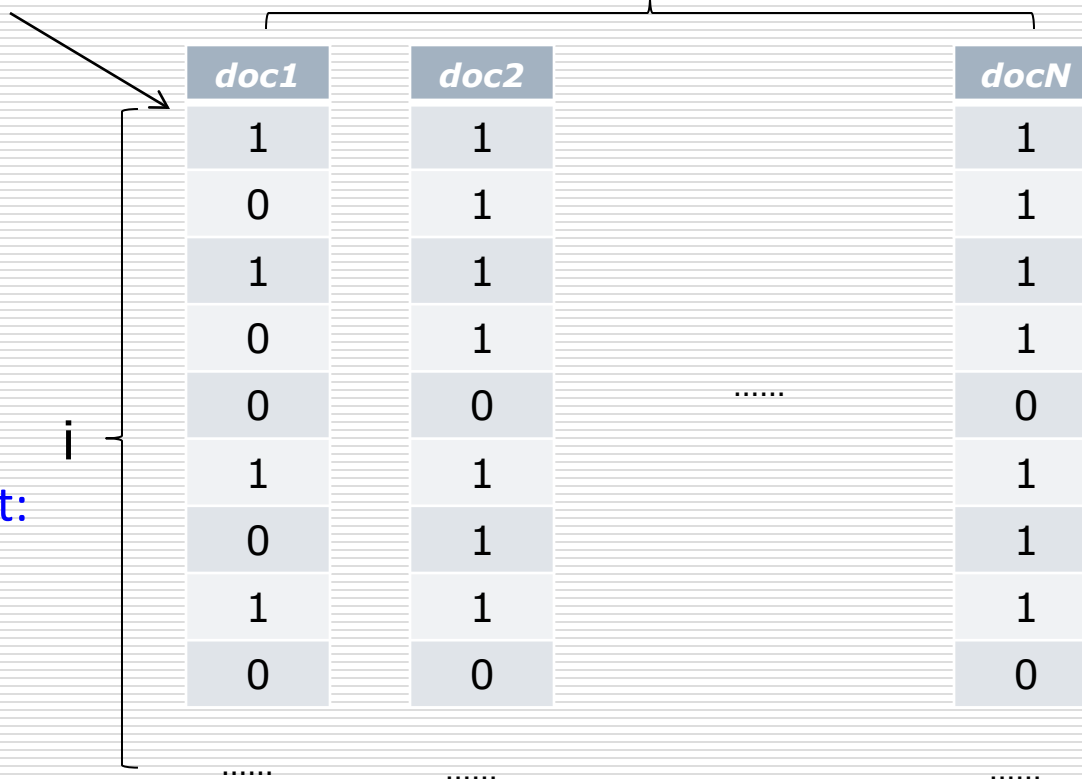
precision? **Shingling**

n

high time complexity:  $O(n^2) \Rightarrow O(n)$

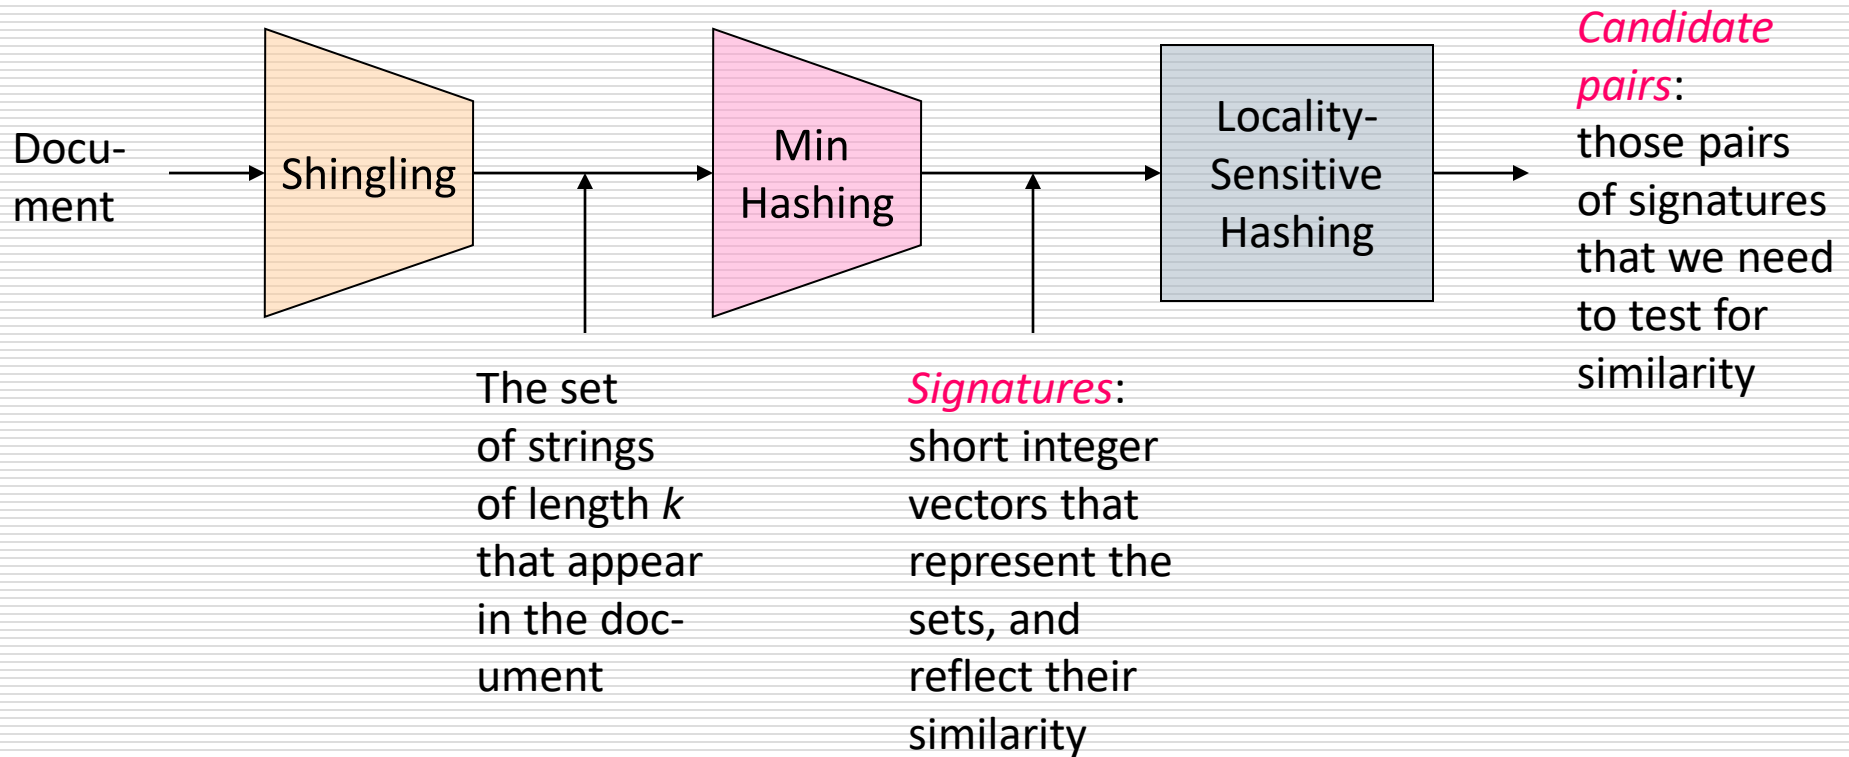
**LSH**

large space  
requirement:  
**MinHash**



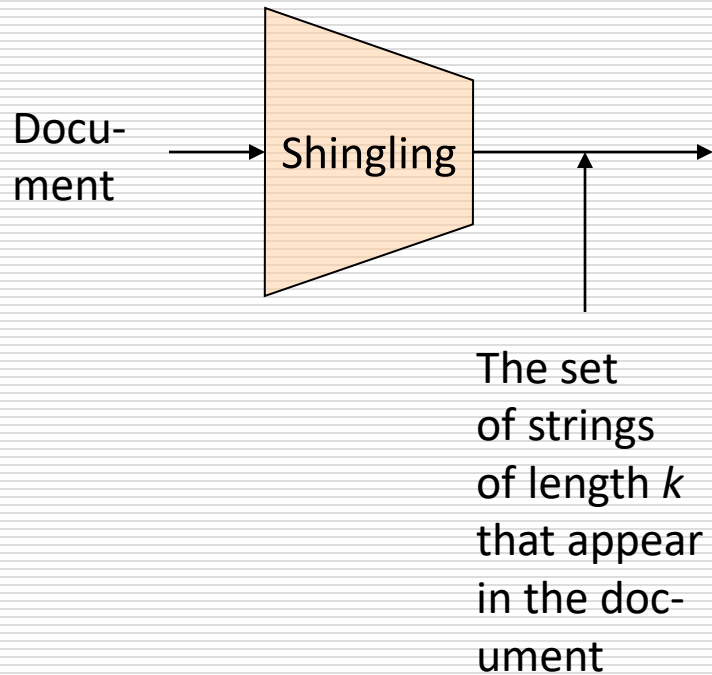
# The Big Picture

---



# The Big Picture

---



□ Step 1: *Shingling*: Convert documents to sets

# Documents as High-Dim Data

---

- Step 1: *Shingling*: Convert documents to sets
- Simple approaches:
  - Document = set of words appearing in document
  - Document = set of “important” words
  - Don’t work well for this application. *Why?*
- Need to account for ordering of words!
- A different way: *Shingles*!



# Define: Shingles (Grams)

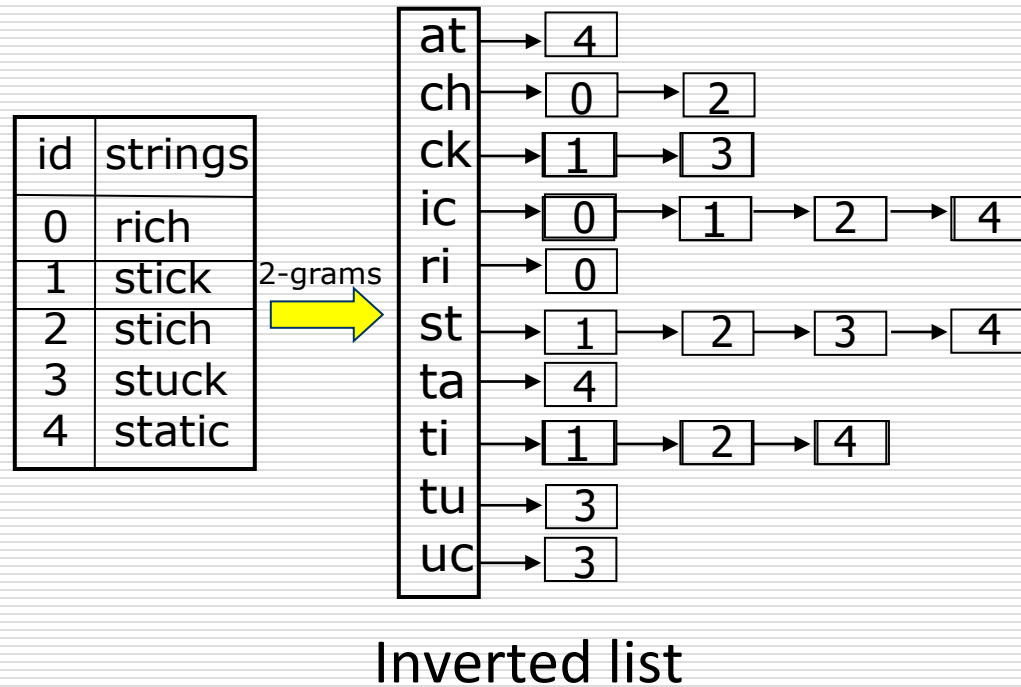
---

- A *k*-shingle (or *k*-gram) for a document is a sequence of *k* tokens that appears in the doc
  - Tokens can be *characters*, *words* or something else, depending on the application
  - Assume tokens = characters for examples
  
- *Example*:  $k=2$ ; string  $S_1 = \text{abcb}$   
Set of 2-shingles:  $S(S_1) = \{\text{ab}, \text{bc}, \text{ca}\}$

# Define: Shingles (Grams)

---

- Application: similar string search

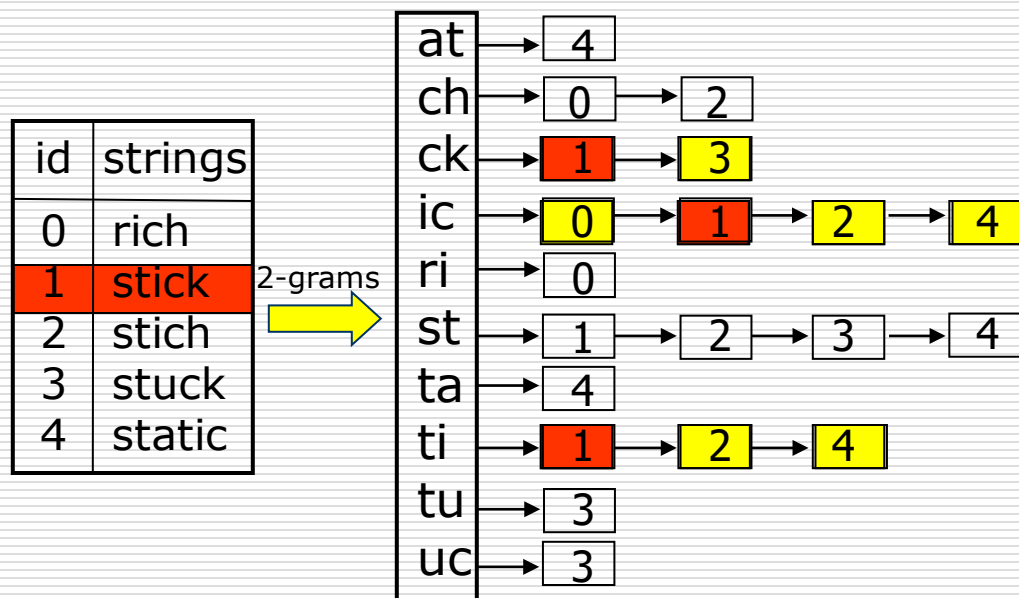


# Define: Shingles (Grams)

- Application: similar string search

Query: "shtick"

sh ht ti ic ck



Inverted list

# Compressing Shingles

---

- To compress long shingles, we can hash them to (say) 4 bytes
- Represent a document by the set of hash values of its  $k$ -shingles
  - Idea: Two documents could (rarely) appear to have shingles in common, when in fact only the hash-values were shared
- Example:  $k=2$ ; document  $D_1 = \text{abcbab}$   
Set of 2-shingles:  $S(D_1) = \{\text{ab}, \text{bc}, \text{ca}\}$   
Hash the singles:  $h(D_1) = \{1, 5, 7\}$

# Similarity Metric for Shingles

---

- Document  $D_1$  is a set of its  $k$ -shingles  $C_1 = S(D_1)$
- Equivalently, each document is a 0/1 vector in the space of  $k$ -shingles
  - Each unique shingle is a dimension
  - Vectors are very sparse

	<i>doc1</i>	<i>doc2</i>
go out	1	1
like travel	0	1
Every day	1	1

.....

.....



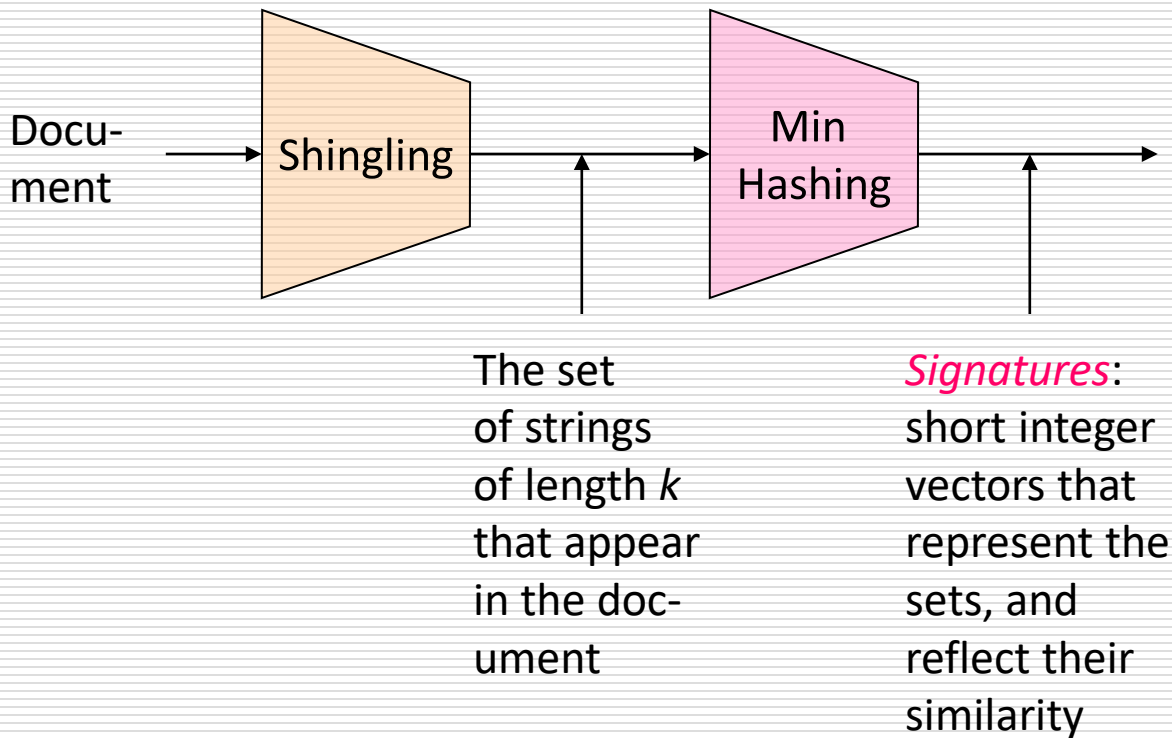
# Working Assumption

---

- Documents that have lots of shingles in common have similar text, even if the text appears in different order
- **Caveat:** You must pick  $k$  large enough, or most documents will have most shingles
  - $k = 5$  is OK for short documents
  - $k = 10$  is better for long documents

# The Big Picture

---



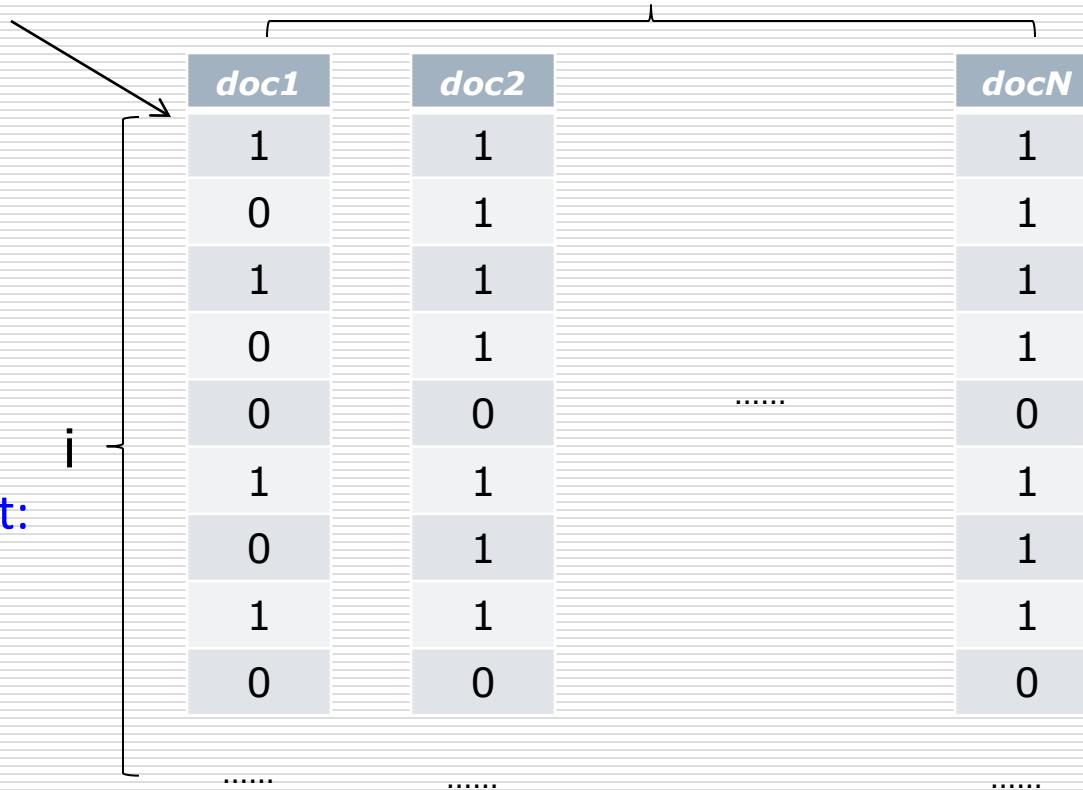
- Step 2: *Minhashing*: Convert large sets to short signatures, while preserving similarity

# Problem for Today's Lecture

---

precision? Shingling

large space  
requirement:  
MinHash



# From Sets to Boolean Matrices

---

- Rows = elements (shingles)
- Columns = sets (documents)
  - 1 in row  $e$  and column  $s$  if and only if  $e$  is a member of  $s$
  - Column similarity is the Jaccard similarity of the corresponding sets (rows with value 1)
  - Typical matrix is sparse!
- Each document is a column:
  - Example:  $\text{sim}(C_1, C_2) = ?$ 
    - Size of intersection = 3; size of union = 6, Jaccard similarity (not distance) =  $3/6$
    - $d(C_1, C_2) = 1 - (\text{Jaccard similarity}) = 3/6$

		Documents			
Shingles		1	1	1	0
		1	1	0	1
		0	1	0	1
		0	0	0	1
		1	0	0	1
		1	1	1	0
		1	0	1	0
		$C_1$	$C_2$		

# Outline: Finding Similar Columns

---

## □ So far:

- Documents → Sets of shingles
- Represent sets as boolean vectors in a matrix

## □ Next goal: Find similar columns while computing small signatures

- Similarity of columns == similarity of signatures



# Outline: Finding Similar Columns

---

- Next Goal: Find similar columns, Small signatures
- Naïve approach:
  - 1) Signatures of columns: small summaries of columns
  - 2) Examine pairs of signatures to find similar columns
    - Essential: Similarities of signatures and columns are related
  - 3) Optional: Check that columns with similar signatures are really similar

# Hashing Columns (Signatures)

---

- **Key idea:** “hash” each column  $C$  to a small *signature*  $h(C)$ , such that:
  - (1)  $h(C)$  is small enough that the signature fits in RAM
  - (2)  $\text{sim}(C_1, C_2)$  is the same as the “similarity” of signatures  $h(C_1)$  and  $h(C_2)$
- **Goal:** Find a hash function  $h(\cdot)$  such that:
  - If  $\text{sim}(C_1, C_2)$  is high, then with high prob.  $h(C_1) = h(C_2)$
  - If  $\text{sim}(C_1, C_2)$  is low, then with high prob.  $h(C_1) \neq h(C_2)$

# Min-Hashing

---

- Goal: Find a hash function  $h(\cdot)$  such that:
  - if  $\text{sim}(C_1, C_2)$  is high, then with high prob.  $h(C_1) = h(C_2)$
  - if  $\text{sim}(C_1, C_2)$  is low, then with high prob.  $h(C_1) \neq h(C_2)$
- Clearly, the hash function depends on the similarity metric:
  - Not all similarity metrics have a suitable hash function
- There is a suitable hash function for the Jaccard similarity: It is called Min-Hashing

# Min-Hashing

---

- Imagine the rows of the boolean matrix permuted under **random permutation**  $\pi$
- Define a **“hash” function**  $h_{\pi}(C)$  = the index of the first (in the permuted order  $\pi$ ) row in which column  $C$  has value 1:

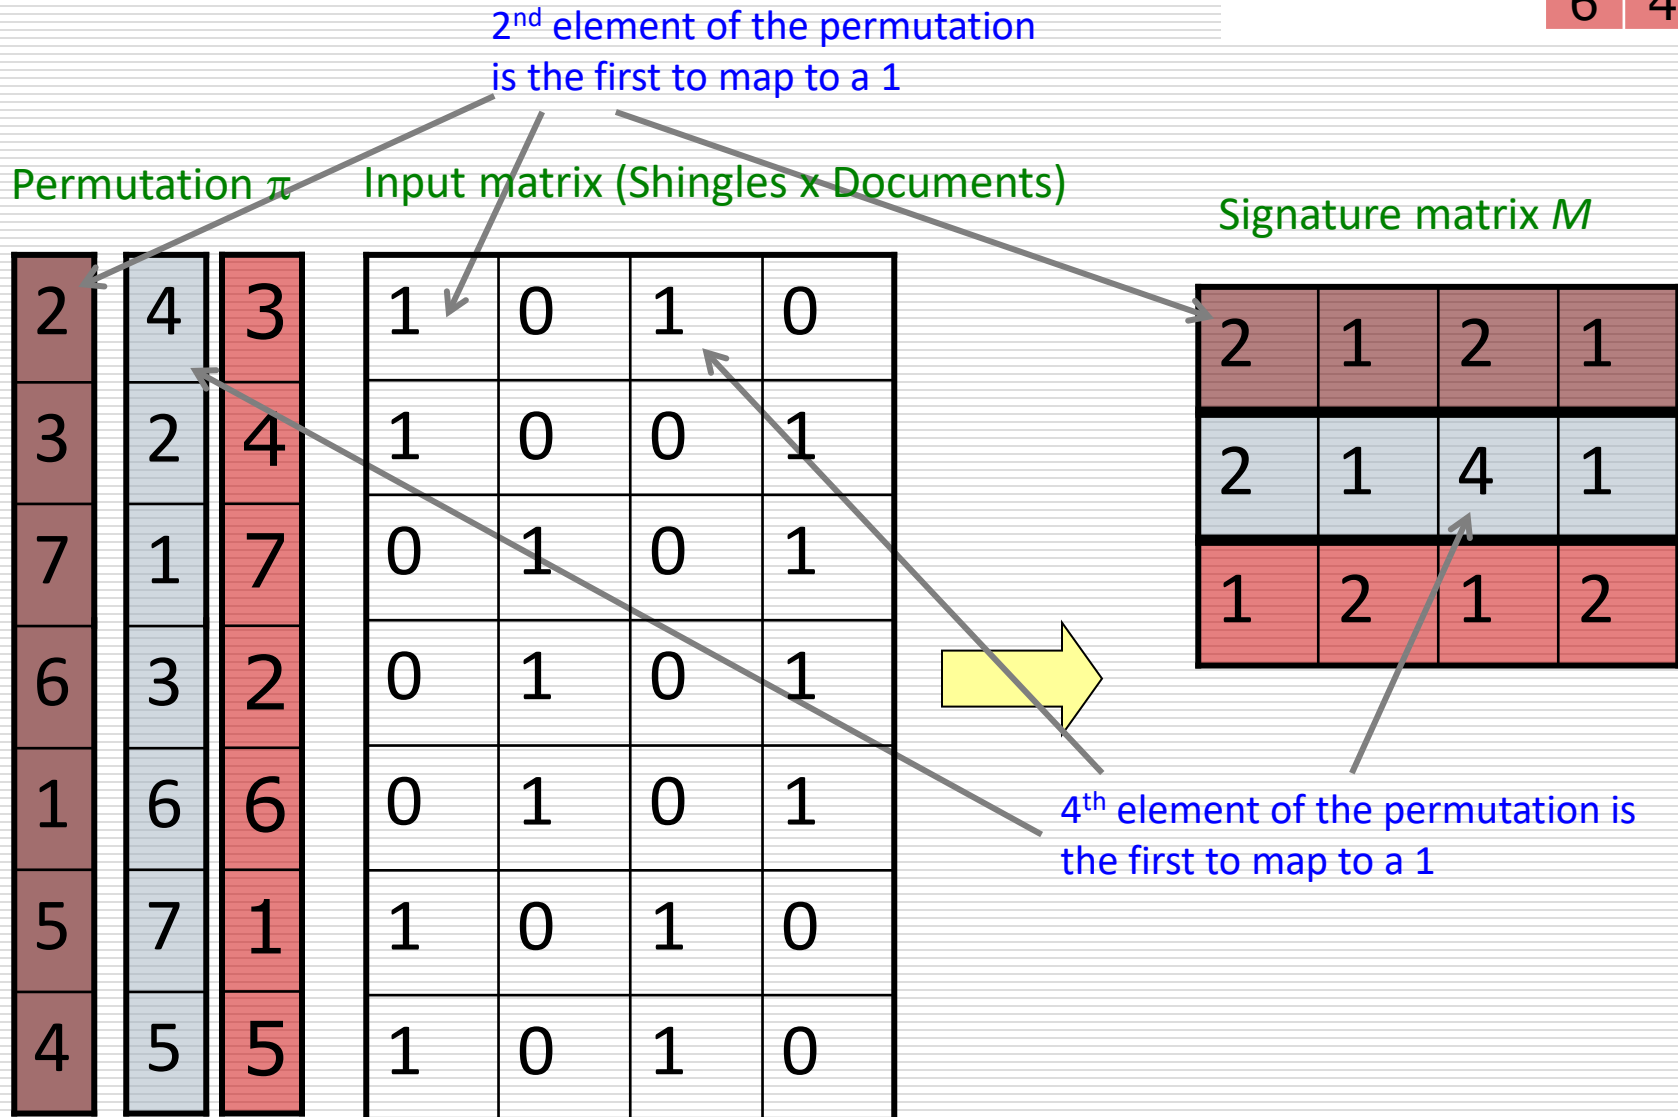
$$h_{\pi}(C) = \min_{\pi} \pi(C)$$

- Use several (e.g., 100) independent hash functions (that is, permutations) to create a signature of a column

# Min-Hashing Example

Note: Another (equivalent) way is to store row indexes:

1	5	1	5
2	3	1	3
6	4	6	4



# Four Types of Rows

---

- Given cols  $C_1$  and  $C_2$ , rows may be classified as:

	$C_1$	$C_2$
A	1	1
B	1	0
C	0	1
D	0	0

- $a$  = # rows of type A, etc.
- Note:  $\text{sim}(C_1, C_2) = a/(a+b+c)$
- Then:  $\Pr[h(C_1) = h(C_2)] = \text{Sim}(C_1, C_2)$ 
  - Look down the cols  $C_1$  and  $C_2$  until we see a 1
  - If it's a type-A row, then  $h(C_1) = h(C_2)$   
If a type-B or type-C row, then not

# The Min-Hash Property

---

- Choose a random permutation  $\pi$
- Claim:  $\Pr[h_\pi(C_1) = h_\pi(C_2)] = \text{sim}(C_1, C_2)$
- Why?

0	0
0	0
1	1
0	0
0	1
1	0

One of the two  
cols had to have  
1 at position  $y$

- Let  $X$  be a doc (set of shingles),  $y \in X$  is a shingle
- Then:  $\Pr[\pi(y) = \min(\pi(X))] = 1/|X|$ 
  - It is equally likely that any  $y \in X$  is mapped to the *min* element
- Let  $y$  be s.t.  $\pi(y) = \min(\pi(C_1 \cup C_2))$
- Then either:  $\pi(y) = \min(\pi(C_1))$  if  $y \in C_1$ , or  
 $\pi(y) = \min(\pi(C_2))$  if  $y \in C_2$
- So the prob. that both are true is the prob.  $y \in C_1 \cap C_2$
- $\Pr[\min(\pi(C_1)) = \min(\pi(C_2))] = |C_1 \cap C_2| / |C_1 \cup C_2| = \text{sim}(C_1, C_2)$

# Similarity for Signatures

---

- We know:  $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
- Now generalize to multiple hash functions
- The *similarity of two signatures* is the fraction of the hash functions in which they agree
- **Note:** Because of the Min-Hash property, the similarity of columns is the same as the expected similarity of their signatures



# Min-Hashing Example

Permutation  $\pi$

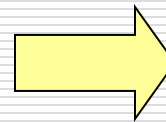
2	4	3
3	2	4
7	1	7
6	3	2
1	6	6
5	7	1
4	5	5

Input matrix (Shingles x Documents)

1	0	1	0
1	0	0	1
0	1	0	1
0	1	0	1
0	1	0	1
1	0	1	0
1	0	1	0

Signature matrix  $M$

2	1	2	1
2	1	4	1
1	2	1	2



Similarities:

	1-3	2-4	1-2	3-4
Col/Col	0.75	0.75	0	0
Sig/Sig	0.67	1.00	0	0

# Min-Hash Signatures

---

- Pick  $K=100$  random permutations of the rows
- Think of  $sig(C)$  as a column vector
- $sig(C)[i]$  = according to the  $i$ -th permutation, the index of the first row that has a 1 in column  $C$   
$$sig(C)[i] = \min(\pi_i(C))$$
- **Note:** The sketch (signature) of document  $C$  is small  
~100 bytes!
- **We achieved our goal!** We “compressed” long bit vectors into short signatures

# Implementation Trick

---

- ❑ **Permuting rows even once is prohibitive**
- ❑ **Row hashing!**
  - Pick  $K = 100$  hash functions  $k_i$
  - Ordering under  $k_i$  gives a random row permutation!
- ❑ **One-pass implementation**
  - For each column  $C$  and hash-func.  $k_i$  keep a “slot” for the min-hash value
  - Initialize all  $sig(C)[i] = \infty$
  - Scan rows looking for 1s
    - ❑ Suppose row  $j$  has 1 in column  $C$
    - ❑ Then for each  $k_i$ :
      - If  $k_i(j) < sig(C)[i]$ , then  $sig(C)[i] \leftarrow k_i(j)$

**How to pick a random hash function  $h(x)$ ?**

**Universal hashing:**

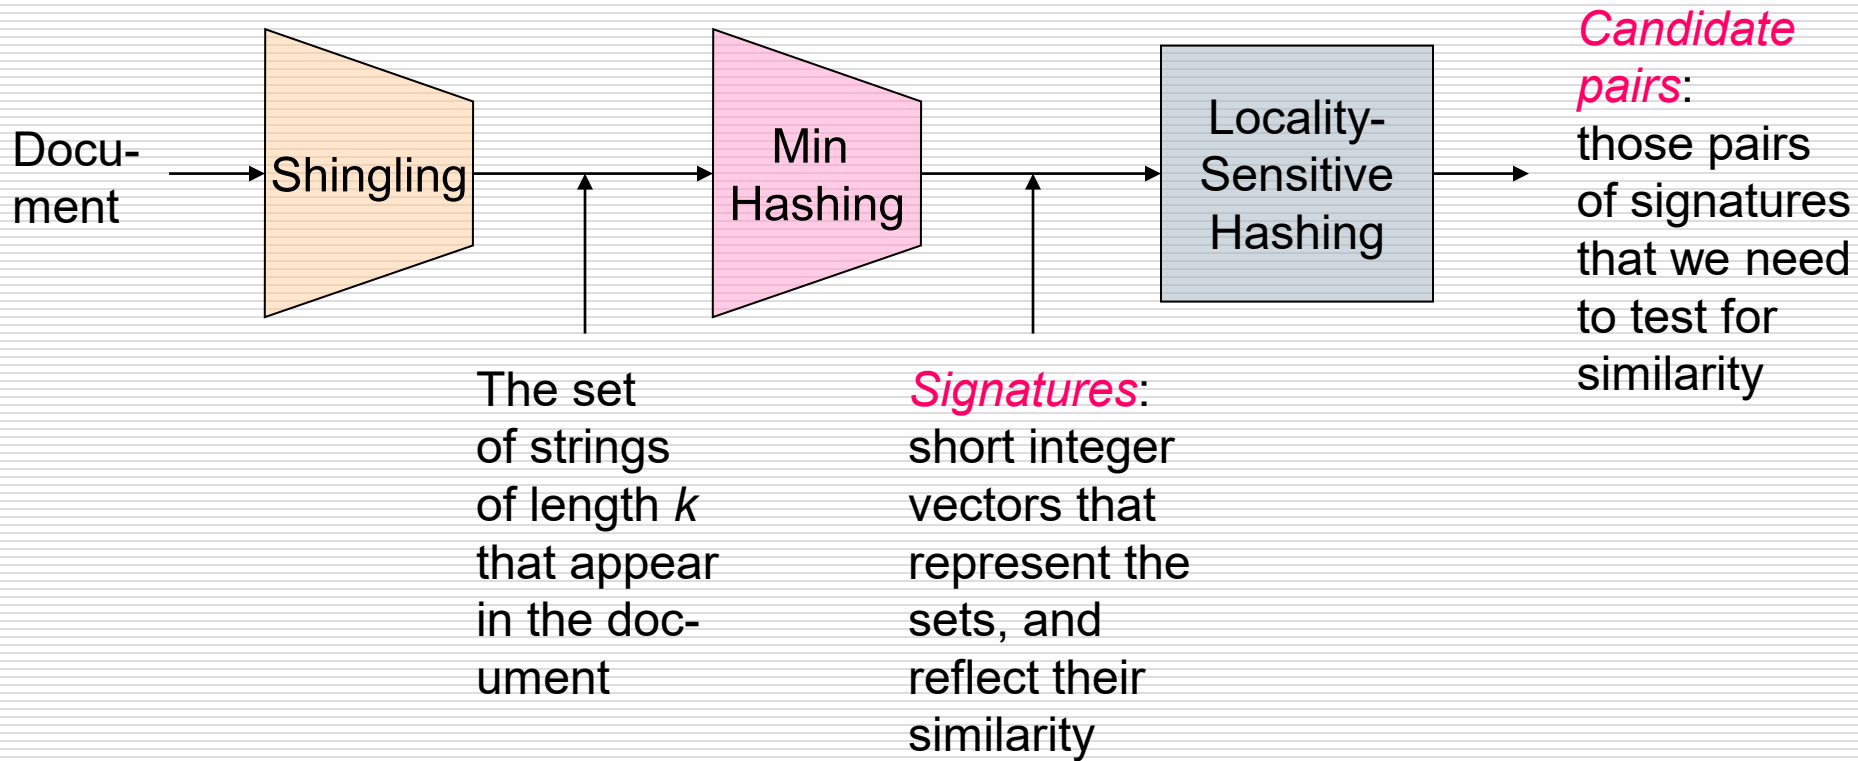
$h_{a,b}(x) = ((a \cdot x + b) \bmod p) \bmod N$   
where:

$a, b \dots$  random integers

$p \dots$  prime number ( $p > N$ )

# The Big Picture

---



- Step 3: *Locality-Sensitive Hashing:*  
Focus on pairs of signatures likely to be from similar documents

# Motivation for LSH

---

- Suppose we need to find near-duplicate documents among  $N = 1$  million documents
- Naïvely, we would have to compute pairwise Jaccard similarities for every pair of docs
  - $N(N - 1)/2 \approx 5 \cdot 10^{11}$  comparisons
  - At  $10^5$  secs/day and  $10^6$  comparisons/sec, it would take 5 days
- For  $N = 10$  million, it takes more than a year...

# Problem for Today's Lecture

---

precision? Shingling

n high time complexity:  $O(n^2) \Rightarrow O(n)$

LSH

large space  
requirement:  
MinHash

doc1	doc2	.....	docN
1	1		1
0	1		1
1	1		1
0	1		1
0	0	.....	0
1	1		1
0	1		1
1	1		1
0	0		0
.....	.....		.....

# LSH: First Cut

---

2	1	4	1
1	2	1	2
2	1	2	1

- Goal: Find documents with Jaccard similarity at least  $s$  (for some similarity threshold, e.g.,  $s=0.8$ )
- LSH – General idea: Use a function  $f(x,y)$  that tells whether  $x$  and  $y$  is a *candidate pair*: a pair of elements whose similarity must be evaluated
- For Min-Hash matrices:
  - Hash columns of *signature matrix*  $M$  to many buckets
  - Each pair of documents that hashes into the same bucket is a *candidate pair*

# Candidates from Min-Hash

---

2	1	4	1
1	2	1	2
2	1	2	1

- Pick a similarity threshold  $s$  ( $0 < s < 1$ )
- Columns  $x$  and  $y$  of  $M$  are a **candidate pair** if their signatures agree on at least fraction  $s$  of their rows:  
 $M(i, x) = M(i, y)$  for at least frac.  $s$  values of  $i$ 
  - We expect documents  $x$  and  $y$  to have the same (Jaccard) similarity as their signatures



# LSH for Min-Hash

---

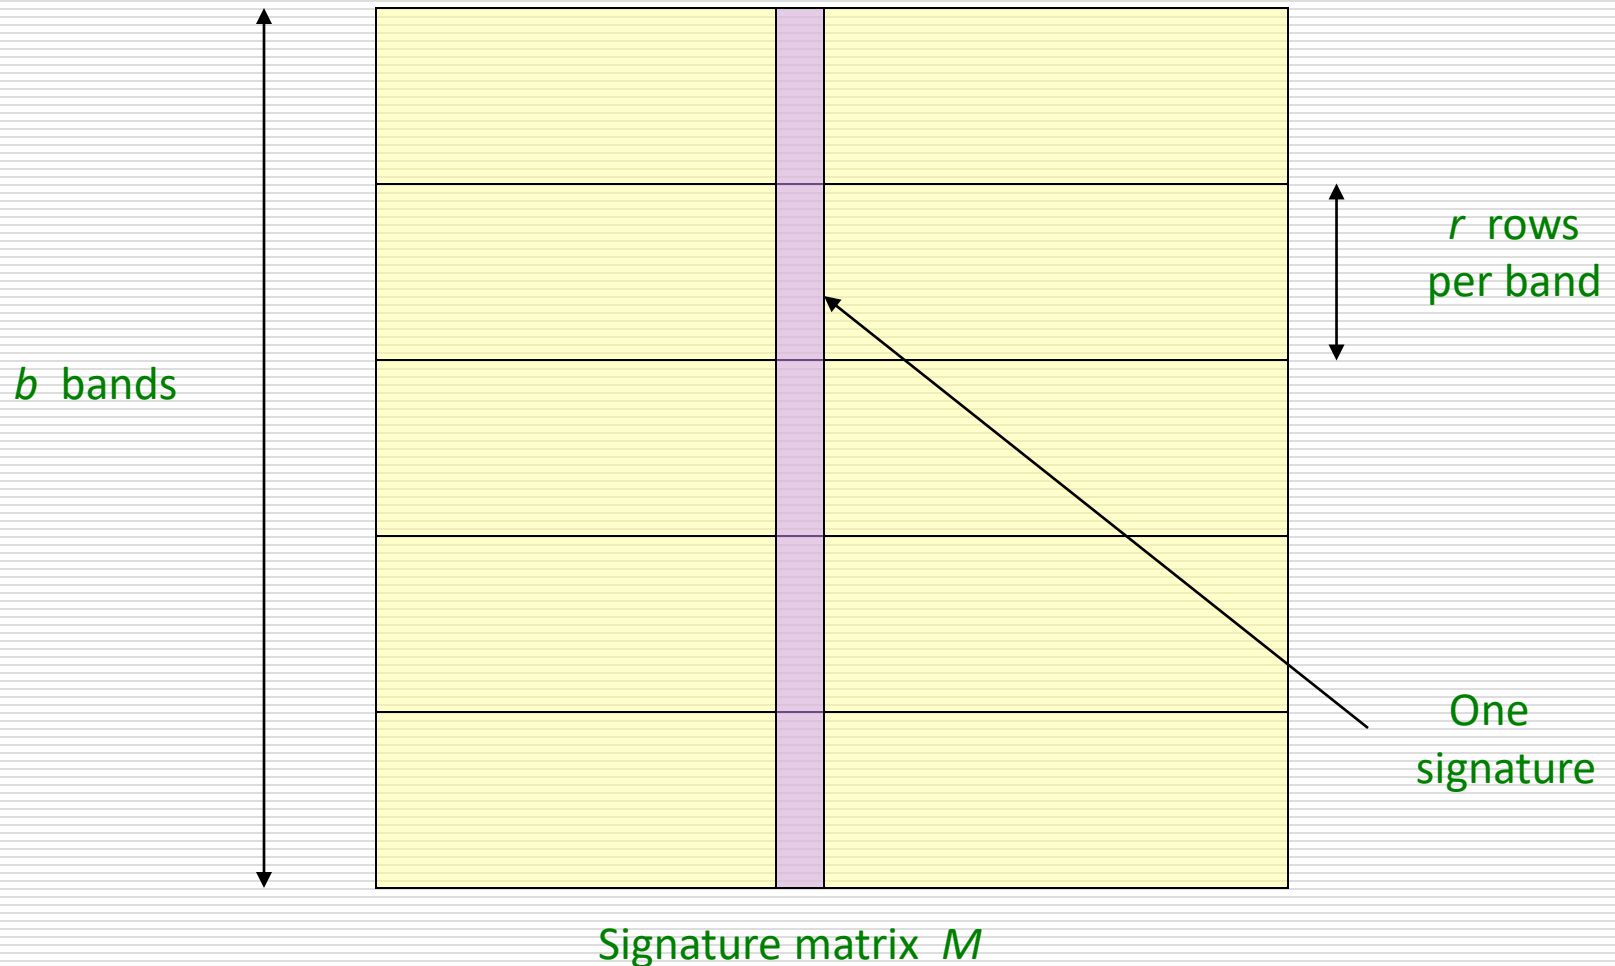
2	1	4	1
1	2	1	2
2	1	2	1

- Big idea: Hash columns of signature matrix  $M$  several times
- Arrange that (only) similar columns are likely to hash to the same bucket, with high probability
- Candidate pairs are those that hash to the same bucket

# Partition $M$ into $b$ Bands

---

2	1	4	1
1	2	1	2
2	1	2	1

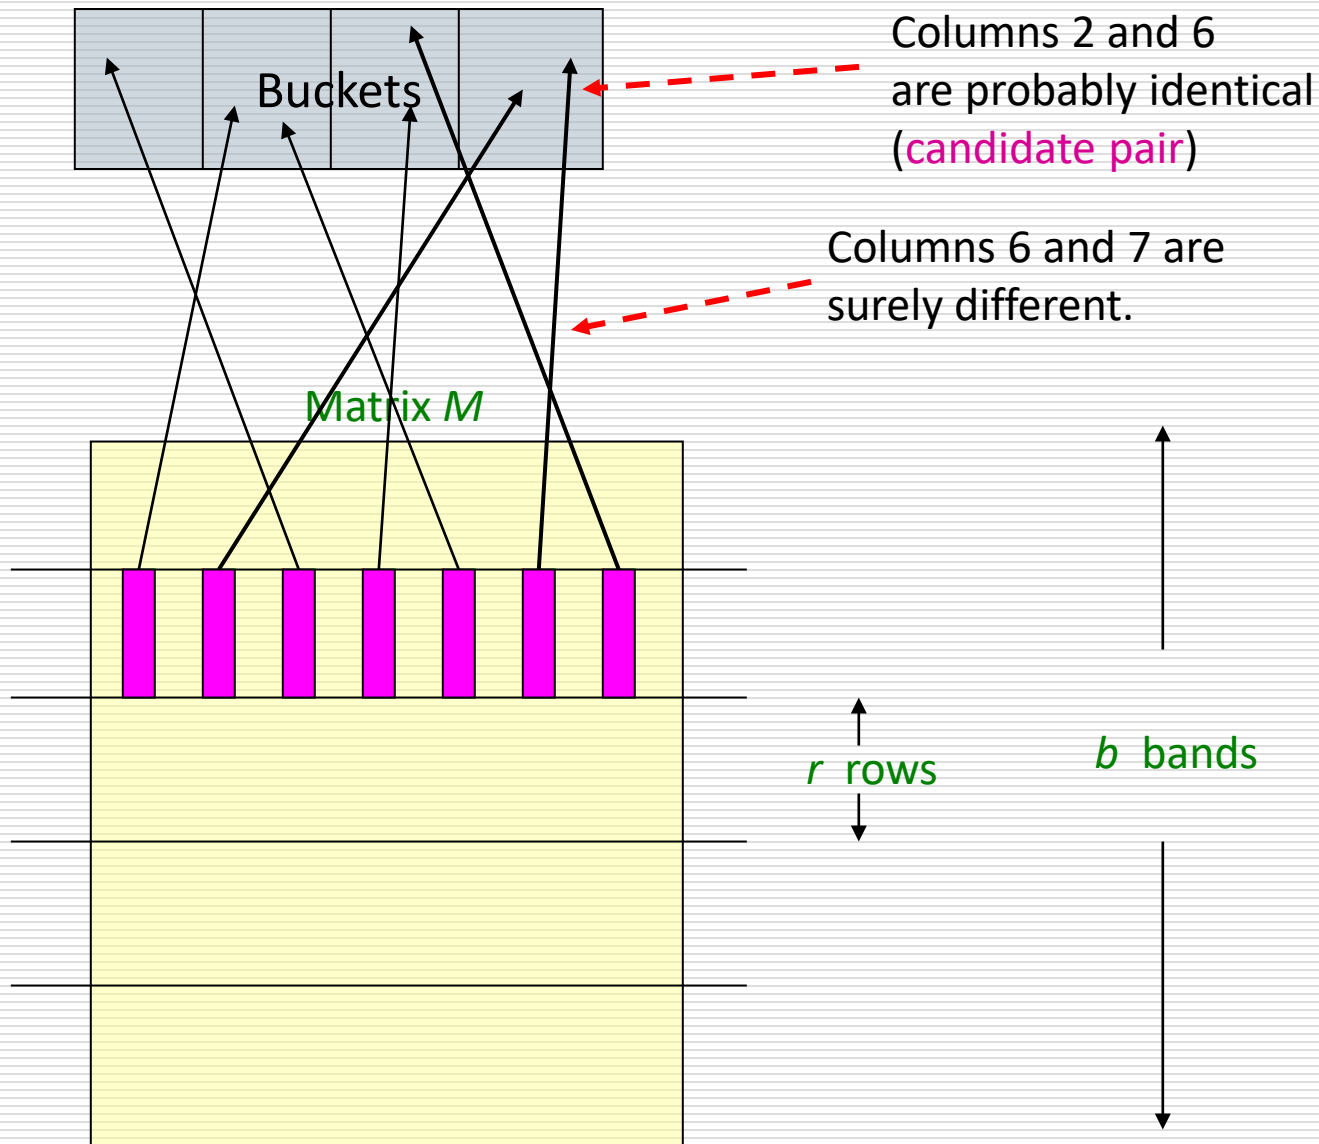


# Partition $M$ into Bands

---

- Divide matrix  $M$  into  $b$  bands of  $r$  rows
- For each band, hash its portion of each column to a hash table with  $k$  buckets
  - Make  $k$  as large as possible
- *Candidate* column pairs are those that hash to the same bucket for  $\geq 1$  band
- Tune  $b$  and  $r$  to catch most similar pairs, but few non-similar pairs

# Hashing Bands



# Simplifying Assumption

---

- There are enough buckets that columns are unlikely to hash to the same bucket unless they are **identical** in a particular band
- Hereafter, we assume that “same bucket” means “identical in that band”
- Assumption needed only to simplify analysis, not for correctness of algorithm

# Example of Bands

---

2	1	4	1
1	2	1	2
2	1	2	1

Assume the following case:

- Suppose 100,000 columns of  $M$  (100k docs)
- Signatures of 100 integers (rows)
- Therefore, signatures take 40Mb
- Choose  $b = 20$  bands of  $r = 5$  integers/band
- Goal: Find pairs of documents that are at least  $s = 0.8$  similar

# $C_1, C_2$ are 80% Similar

2	1	4	1
1	2	1	2
2	1	2	1

- Find pairs of  $\geq s=0.8$  similarity, set  $b=20$ ,  $r=5$
- Assume:  $\text{sim}(C_1, C_2) = 0.8$ 
  - Since  $\text{sim}(C_1, C_2) \geq s$ , we want  $C_1, C_2$  to be a **candidate pair**:  
We want them to hash to at **least 1 common bucket** (at least one band is identical)
- Probability  $C_1, C_2$  identical in one particular band:  $(0.8)^5 = 0.328$
- Probability  $C_1, C_2$  are **not** similar in all of the 20 bands:  $(1-0.328)^{20} = 0.00035$ 
  - i.e., about 1/3000th of the 80%-similar column pairs are **false negatives** (we miss them)
  - We would find 99.965% pairs of truly similar documents

# $C_1, C_2$ are 30% Similar

2	1	4	1
1	2	1	2
2	1	2	1

- Find pairs of  $\geq s=0.8$  similarity, set  $b=20$ ,  $r=5$
- Assume:  $\text{sim}(C_1, C_2) = 0.3$ 
  - Since  $\text{sim}(C_1, C_2) < s$  we want  $C_1, C_2$  to hash to **NO common buckets** (all bands should be different)
- Probability  $C_1, C_2$  identical in one particular band:  
 $(0.3)^5 = 0.00243$
- Probability  $C_1, C_2$  identical in at least 1 of 20 bands:  
 $1 - (1 - 0.00243)^{20} = 0.0474$ 
  - In other words, approximately 4.74% pairs of docs with similarity 0.3% end up becoming **candidate pairs**
    - They are **false positives** since we will have to examine them (they are candidate pairs) but then it will turn out their similarity is below threshold  $s$



# LSH Involves a Tradeoff

---

2	1	4	1
1	2	1	2
2	1	2	1

## □ Pick:

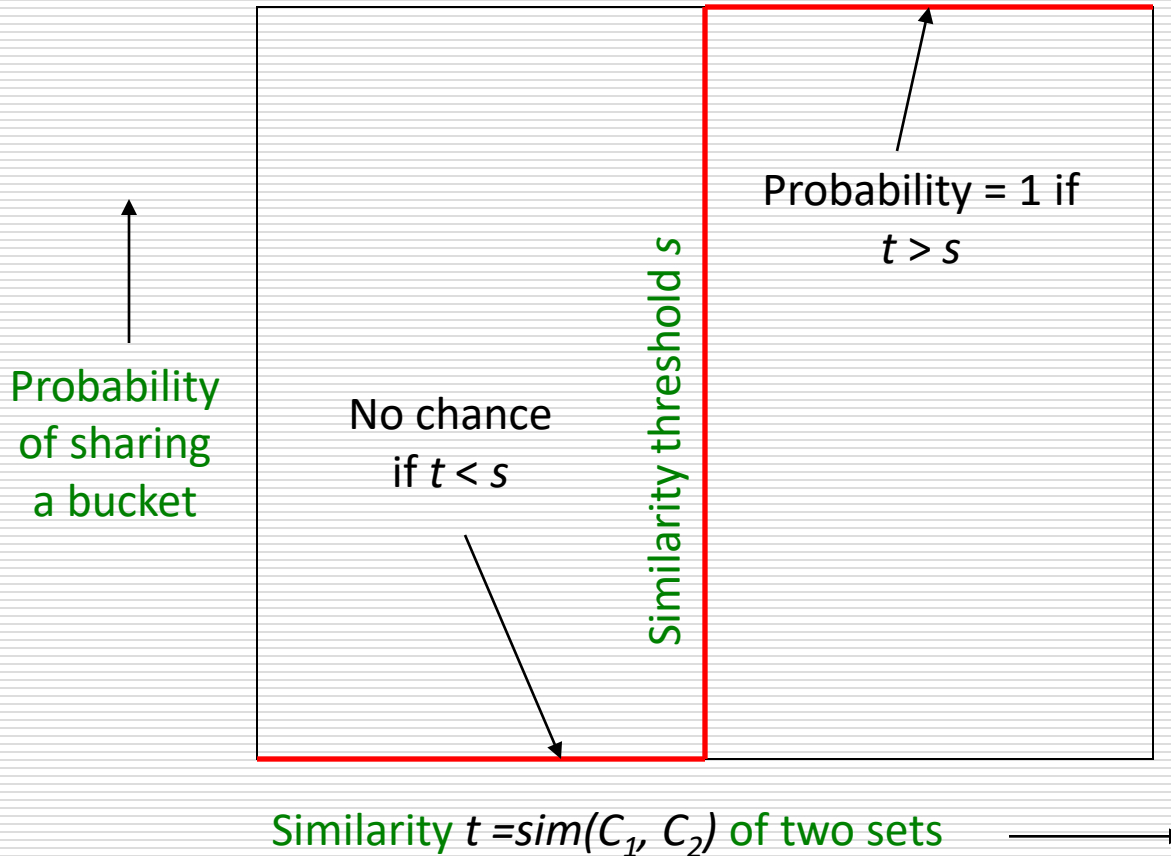
- The number of Min-Hashes (rows of  $M$ )
- The number of bands  $b$ , and
- The number of rows  $r$  per band

to balance false positives/negatives

- **Example:** If we had only 15 bands of 5 rows, the number of false positives would go down, but the number of false negatives would go up

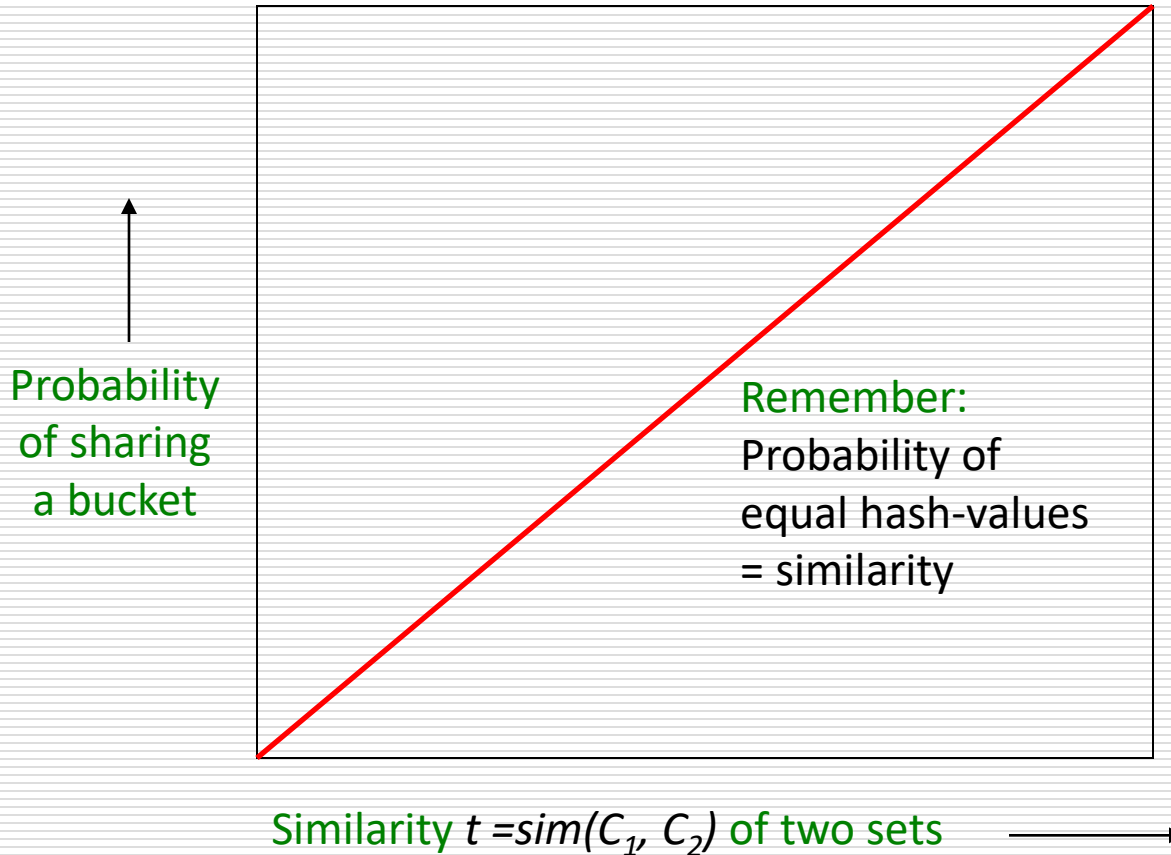
# Analysis of LSH – What We Want

---



# What 1 Band of 1 Row Gives You

---

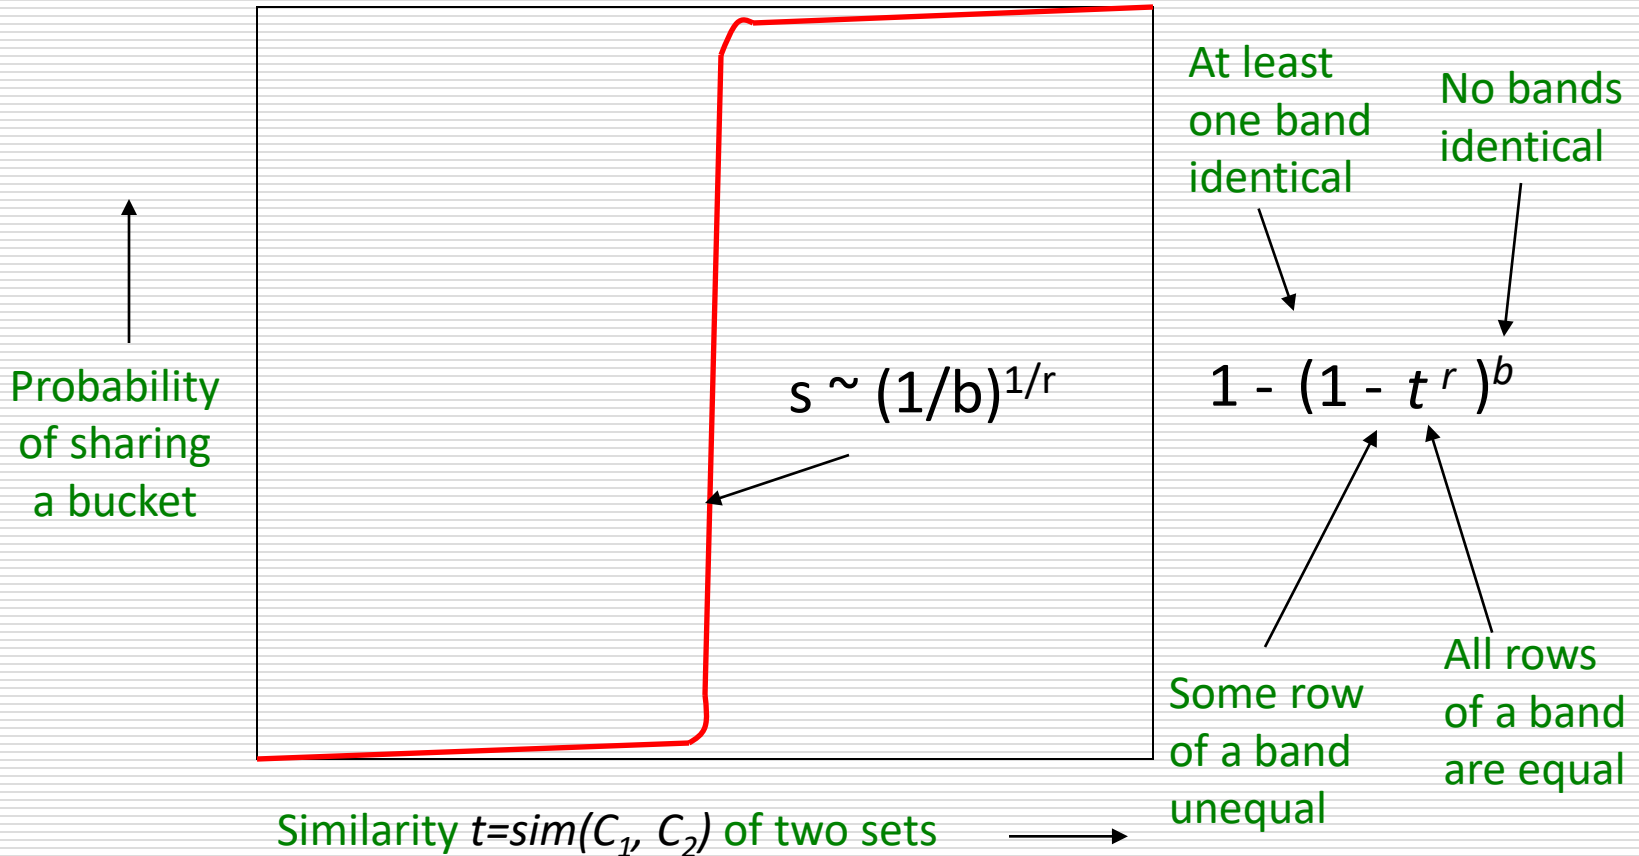


## $b$ bands, $r$ rows/band

---

- Columns  $C_1$  and  $C_2$  have similarity  $t$
- Pick any band ( $r$  rows)
  - Prob. that all rows in band equal =  $t^r$
  - Prob. that some row in band unequal =  $1 - t^r$
- Prob. that no band identical =  $(1 - t^r)^b$
- Prob. that at least 1 band identical =  $1 - (1 - t^r)^b$

# What $b$ Bands of $r$ Rows Gives You



Example:  $b = 20; r = 5$

---

□ **Similarity threshold  $s$**

□ **Prob. that at least 1 band is identical:**

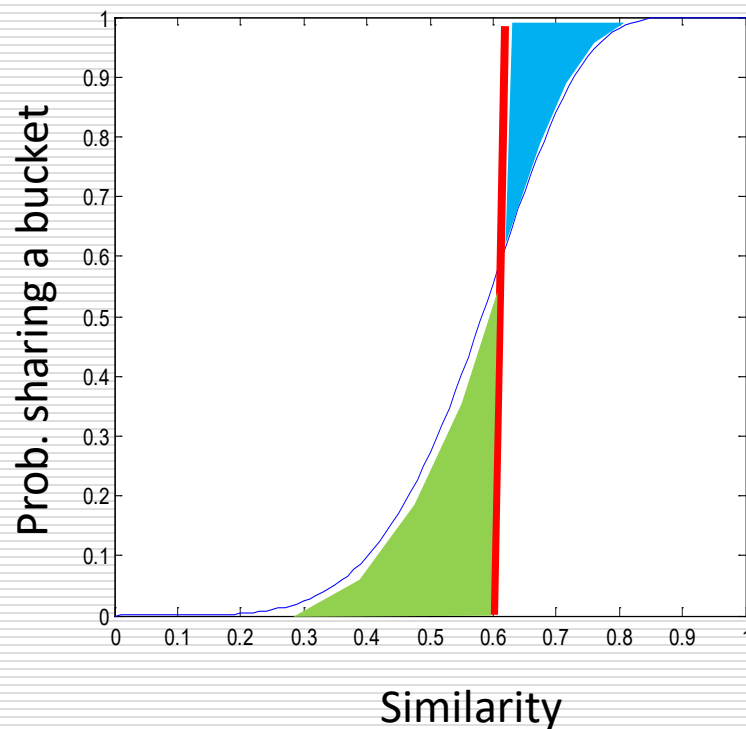
$s$	$1-(1-s^r)^b$
.2	.006
.3	.047
.4	.186
.5	.470
.6	.802
.7	.975
.8	.9996

# Picking $r$ and $b$ : The S-curve

---

□ Picking  $r$  and  $b$  to get the best S-curve

■ 50 hash-functions ( $r=5$ ,  $b=10$ )



Blue area: False Negative rate

Green area: False Positive rate

# LSH Summary

---

- Tune  $M, b, r$  to get almost all pairs with similar signatures, but eliminate most pairs that do not have similar signatures
- Check in main memory that **candidate pairs** really do have **similar signatures**
- **Optional:** In another pass through data, check that the remaining candidate pairs really represent similar documents



# Summary: 3 Steps

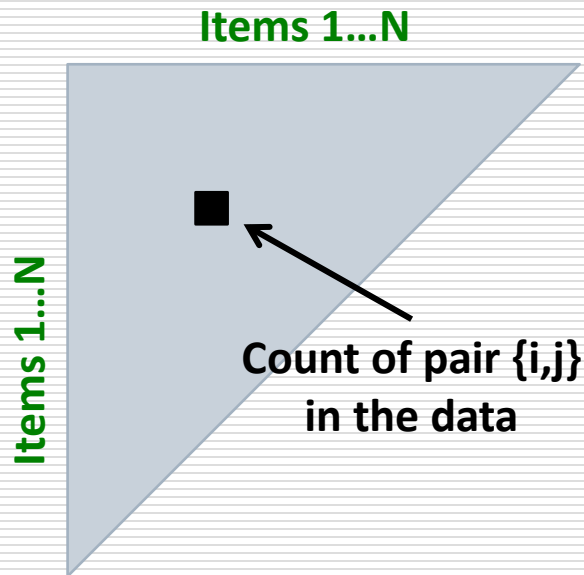
---

- **Shingling:** Convert documents to sets
  - We used hashing to assign each shingle an ID
- **Min-Hashing:** Convert large sets to short signatures, while preserving similarity
  - We used **similarity preserving hashing** to generate signatures with property  $\Pr[h_{\pi}(C_1) = h_{\pi}(C_2)] = \text{sim}(C_1, C_2)$
  - We used hashing to get around generating random permutations
- **Locality-Sensitive Hashing:** Focus on pairs of signatures likely to be from similar documents
  - We used hashing to find **candidate pairs** of similarity  $\geq s$

# Relation to Previous Lecture

---

## □ Last time: Finding frequent pairs

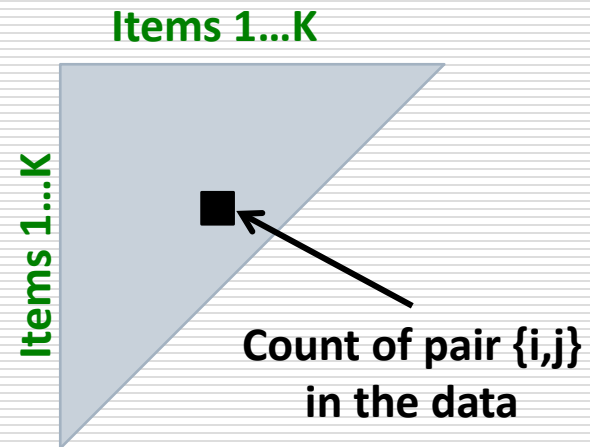


### Naïve solution:

Single pass but requires space quadratic in the number of items

N ... number of distinct items

K ... number of items with support  $\geq s$



### A-Priori:

First pass: Find frequent singletons

For a pair to be a frequent pair candidate, its singletons have to be frequent!

Second pass:

Count only candidate pairs!

# Relation to Previous Lecture

---

□ Last time: Finding frequent pairs

□ Further improvement: PCY

■ Pass 1:

□ Count exact frequency of each item:



□ Take pairs of items  $\{i,j\}$ , hash them into  $B$  buckets and count of the number of pairs that hashed to each bucket:



Basket 1:  ~~$\{1,2,3\}$~~

Pairs:  $\{1,2\}$   $\{1,3\}$   $\{2,3\}$

# Relation to Previous Lecture

---

□ Last time: Finding frequent pairs

□ Further improvement: PCY

■ Pass 1:

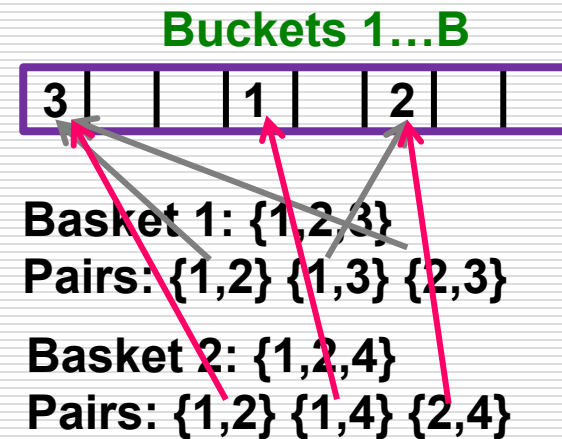
□ Count exact frequency of each item:



□ Take pairs of items  $\{i,j\}$ , hash them into  $B$  buckets and count of the number of pairs that hashed to each bucket:

■ Pass 2:

□ For a pair  $\{i,j\}$  to be a **candidate for a frequent pair**, its singletons  $\{i\}$ ,  $\{j\}$  have to be frequent and the pair has to hash to a frequent bucket!



# Relation to Previous Lecture

---

□ Last time: Finding frequent pairs

□ **Previous lecture: A-Priori**

■ Main idea: Candidates

Instead of keeping a count of each pair, only keep a count of candidate pairs!

**Today's lecture: Find pairs of similar docs**

Main idea: Candidates

■ -- Pass 1: Take documents and hash them to buckets such that documents that are similar hash to the same bucket

-- Pass 2: Only compare documents that are candidates (i.e., they hashed to a same bucket)

Benefits: Instead of  $O(N^2)$  comparisons, we need  $O(N)$  comparisons to find similar documents



3}

{2,4}

# Acknowledgement

---

- Slides are adapted from:
  - Prof. Jeffrey D. Ullman
  - Dr. Anand Rajaraman
  - Dr. Jure Leskovec